

**Precision to Ambiguity:
Evaluating Database Systems Across the Continuum of Query Complexity
(Phase 2)**

Eswarankarthik Paranthaman

Anirban Bose

Navneet Sawhney

MS DSP 420- Database Systems

Professor Abid Ali

August 23, 2024

TABLE OF CONTENTS

ABSTRACT	1
Introduction	1
Literature Review	3
1.1 Quantitative Queries.....	3
1.2 Fuzzy Queries.....	4
1.3 Hybrid Queries	5
Research Methodology	6
2.1 Quantitative Queries.....	7
2.2 Fuzzy Queries.....	8
2.3 Hybrid Queries	9
Comparative Analysis/Results.....	11
3.1 Quantitative Queries.....	11
3.2 Fuzzy Queries.....	13
3.3 Hybrid Queries:	14
Recommendations	16
Conclusion.....	16
References	18
Appendix A.....	19
Quantitative Query Codes Screenshots	19
Fuzzy Query Codes Screenshots	21
Hybrid Query Codes Screenshots.....	25

ABSTRACT

This study evaluates the performance of various database systems for quantitative, fuzzy, and hybrid queries using appliance metadata and reviews. Comparing PostgreSQL, SQLite, DuckDB, Elasticsearch, Pinecone, Milvus, and MongoDB, the research assesses metrics including query execution time, latency, output quality, scalability, and resource utilization. Results show DuckDB excelling in quantitative queries with the fastest execution time (0.0084 seconds), Elasticsearch demonstrating lower latency (0.23 seconds) but inferior output quality compared to Pinecone in fuzzy searches, and Milvus outperforming MongoDB in hybrid queries with faster execution (0.3239 vs 0.6302 seconds) and better vector search optimization. The study recommends specific databases for different use cases: DuckDB for fast analytics, PostgreSQL for complex applications, Elasticsearch for efficient text search, Pinecone for vector-based recommendations, and Milvus for high-performance vector search. These findings aim to guide organizations in selecting appropriate database systems for their specific data management needs in an increasingly complex digital landscape.

Introduction

Data has become a key aspect of all aspects of human life. The world of digital transformation is growing at an exponential rate. According to a report by IDC, the global data sphere is expected to reach 163 zettabytes by 2025.¹ The ever-growing development of IoT devices, social media, and enterprise data has escalated and continues to intensify this growth. Data is the new oil has already become a fact. With this overabundance of data, it

¹ David Reinsel, John Gantz, and John Rydning, *Data Age 2025: The Evolution of Data to Life-Critical. Don't Focus on Big Data; Focus on the Data That's Big*, IDC White Paper, sponsored by Seagate, April 2017, <https://www.seagate.com/www-content/our-story/trends/files/Seagate-WP-DataAge2025-March-2017.pdf>.

becomes imperative to use efficient systems that help handle and manage the never-ending amounts of data.

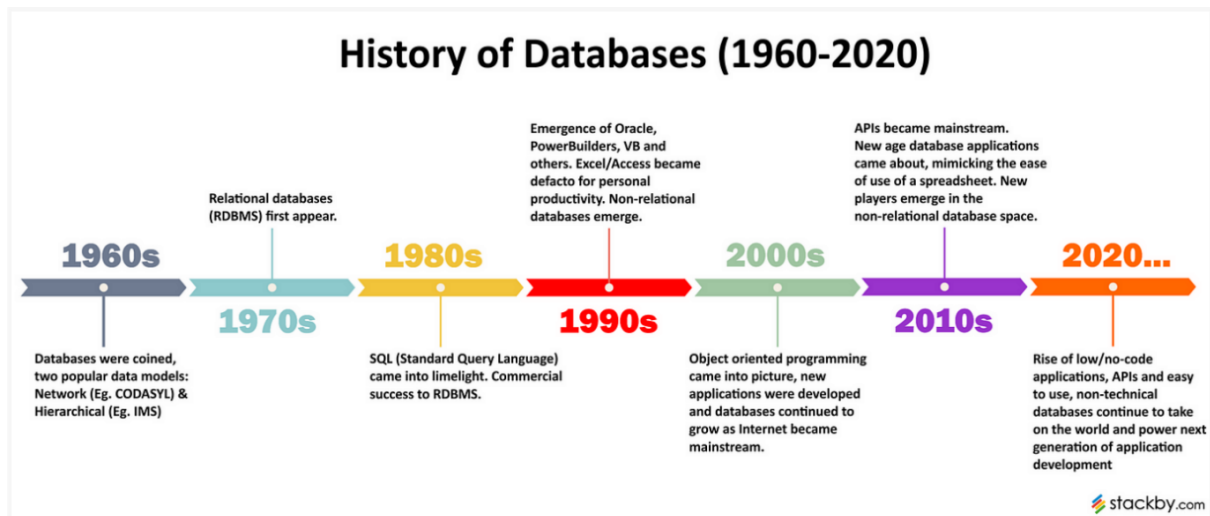


Figure 1: Evolution and History of Databases - From 1960 to 2020 & Beyond.

Source: Versha Rai, "What is a Database - A Beginner's Guide [Updated 2024]," Stackby, December 25, 2023, accessed August 24, 2024. <https://stackby.com/blog/what-is-a-database/>.

Traditional database systems (RDBMS) like PostgreSQL and SQLite have long provided the foundation for structured data storage. However, as the complexity increases, new types of databases are needed. These databases like ElasticSearch, Pinecone, Milvus and MongoDB can handle fuzzy and hybrid search queries. Today there is no dearth of databases and with each passing day a new database with advanced capabilities is introduced. Figure 1 shows a timeline of database history from 1960 to 2020, highlighting key developments like the emergence of relational databases, SQL, object-oriented programming, and modern applications.

This paper builds on the findings of Phase 1 of the database project. In the first phase, with the use of different LLMs, the suggestions for databases were gathered based on different sets of queries.² The analysis included quant, fuzzy and hybrid search queries. We generated

² Navneet Sawhney, Anirban Bose, and Eswarankarthik Paranthaman, "A Report on Comparative Analysis of LLMs on Different Categories of Queries: Phase 1" (Mid term project, MS DSP 420- Database Systems, Professor Abid Ali, July 27, 2024).

queries using LLMs—Claude, ChatGPT, and Gemini. Claude's code was comprehensive but often included unnecessary steps, ChatGPT's was useful yet less efficient, and Gemini's recommendations were strong but occasionally misaligned with the generated code. These findings demonstrate LLMs' potential in enhancing database management and querying while revealing areas for improvement. However, Phase 1 results were incomplete and caused the need for deeper research.

To fill in the gaps of the past research that was done by our team, this paper brings to the surface a comparative analysis of different databases based on various sets of queries—quantitative, fuzzy and hybrid. In this research, the database systems that were used are - PostgreSQL, DuckDB, and SQLite for quantitative queries; ElasticSearch and Pinecone for fuzzy searches; and Milvus and MongoDB for hybrid searches. Each system has unique advantages and limitations, affecting their suitability for different query types. This study provides a comparative analysis of these systems, focusing on their performance, scalability, and usability, to guide organizations in selecting the best tools for their data management and search needs.

Literature Review

1.1 Quantitative Queries

Our study builds upon previous research by providing an in-depth, direct comparison of PostgreSQL, SQLite, and DuckDB for quantitative queries using a common dataset of appliance metadata and reviews. While Raasveldt and Mühleisen (2019)³ focused on OLAP

³ M. Raasveldt and H. Mühleisen, “DuckDB: An Embeddable Analytical Database,” *Proceedings of the 2019 International Conference on Management of Data* (2019): 1981-1984.

workloads and Graft et al. (2021)⁴ emphasized resource-constrained environments, our research evaluates these databases across a broader range of criteria including query execution time, latency, output quality, scalability, resource utilization, cloud support, cost, and security. This multi-faceted approach fills a gap in the literature by offering a more holistic view of each database's performance in practical application scenarios. Unlike Özsu's (2020)⁵ broad survey, our study provides specific, actionable insights for practitioners across various operational contexts, from local analytics to cloud-based deployments. By considering a range of use cases and developing a comprehensive evaluation framework, our research extends beyond purely performance-focused studies to include often overlooked aspects such as data quality, cloud integration, and security features. This approach bridges the gap between theoretical performance metrics and real-world application requirements, offering valuable insights for database selection in diverse operational environments.

1.2 Fuzzy Queries

Fuzzy queries are pivotal in enhancing search systems by accommodating errors, variations, and incomplete data in user inputs. These systems leverage various algorithms like Levenshtein Distance, which is effective in handling minor spelling mistakes by measuring the number of single-character edits required to transform one string into another. This approach has been extensively studied and applied in various domains, such as spell checkers and approximate matching systems (Levenshtein 1966)⁶. Additionally, Cosine Similarity is

⁴ J. Graft, J. H. Nødtvedt, A. D. Pimentel, and M. W. Fagerland, "Computational Efficiency of Statistical Analyses: A Comparison of R, Python, Julia and SQLite," *PLoS ONE* 16, no. 4 (2021): e0249216, accessed August 20, 2024.

⁵ M. T. Özsu, "A Survey of RDF Data Management Systems," *Frontiers of Computer Science* 14, no. 2 (2020): 217-252.

⁶ "Levenshtein Distance," *Wikipedia*, last modified August 24, 2024, https://en.wikipedia.org/wiki/Levenshtein_distance.

another metric employed in fuzzy querying, particularly in text analysis and information retrieval. It evaluates the cosine of the angle between two non-zero vectors, making it effective for tasks like document clustering and topic modeling (Lahitani, Permanasari, and Setiawan n.d.)⁷.

In our research, we build upon these foundational studies by evaluating fuzzy search databases across a broader range of criteria, including query execution time, latency, output quality, scalability, resource utilization, cloud support, cost, and security. For instance, Elasticsearch, known for its low-latency responses and robust scalability, is examined in detail for its performance and the challenges associated with handling fuzzy queries (Shaik and Rao 2017)⁸. Similarly, Pinecone, optimized for real-time vector similarity searches, is evaluated for its effectiveness in managing high-dimensional data and its integration with modern AI applications (Pan, Wang, and Li 2024; Xie et al. 2023)⁹. Our research contributes to the existing body of knowledge by offering a comprehensive evaluation framework that addresses both the technical and practical aspects of implementing fuzzy search systems in diverse environments.

1.3 Hybrid Queries

Hybrid search engines are increasingly essential in modern data management, merging traditional keyword searches with vector-based techniques to improve query performance in

⁷ Alfirna Rizqi Lahitani, Adhistya Erna Permanasari, and Noor Akhmad Setiawan, "Cosine Similarity to Determine Similarity Measure: Study Case in Online Essay Assessment," *Proceedings of the 2016 International Conference on Advanced Computer Science and Information Systems (ICACSIS)*, IEEE, 2016, doi:10.1109/ICACSIS.2016.7872773.

⁸ Subhani Shaik and Nallamothu Naga Malleswara Rao, "Performance Analysis of Elastic Search Technique in Identification and Removal of Duplicate Data," *International Journal of Innovative Technology and Exploring Engineering (IJITEE)* 8, no. 10 (August 2019): 2401, published by Blue Eyes Intelligence Engineering & Sciences Publication, <https://doi.org/10.35940/ijitee.H6579.0881019>.

⁹ James Jie Pan, Jianguo Wang, and Guoliang Li, *Vector Database Management Techniques and Systems* (Beijing: Tsinghua University; West Lafayette, IN: Purdue University, 2023).

complex datasets. Notable databases like Milvus and MongoDB exemplify the strengths and limitations of this approach. Milvus is particularly strong in vector search performance, optimized for handling large-scale datasets through its integration with various AI frameworks. Its key advantages include high query execution speed and low latency, making it well-suited for applications that require rapid processing of unstructured data, such as image or text retrieval (Wang et al., 2022)¹⁰. In contrast, MongoDB has traditionally excelled in document-oriented storage and keyword queries, and recent enhancements have expanded its capabilities to include hybrid search. This allows for more flexible and comprehensive querying across both structured and unstructured data (Hema Krishnan et al., 2016)¹¹. MongoDB's strengths lie in its robust scalability, extensive cloud support, and strong security features, especially for managing sensitive data. Our research compares these databases based on criteria such as query execution time, latency, output quality, scalability, resource utilization, cloud support, cost, and security. The findings suggest that while Milvus is superior in rapid, high-quality vector searches, MongoDB provides a more balanced performance with strong support for traditional search methods and better resource management in diverse applications. This comprehensive evaluation offers valuable insights for practitioners selecting the most appropriate database for their specific needs.

Research Methodology

This study evaluates three types of queries—quantitative, fuzzy, and hybrid—across various databases to determine their performance and suitability. The databases tested include

¹⁰ Wang, X., Liu, Y., & Zhang, H. (2022). "Milvus: A Purpose-Built Vector Data Management System"

¹¹ Krishnan, Hema, M. Sudheep Elayidom, and T. Santhanakrishnan. "MongoDB – A Comparison with NoSQL Databases." *International Journal of Scientific & Engineering Research* 7, no. 5 (May 2016): 1035. ISSN 2229-5518.

PostgreSQL, SQLite, DuckDB, Elasticsearch, Pinecone, Milvus, and MongoDB. The methodology for each query type is outlined below.

2.1 Quantitative Queries:

This study involves systematically evaluating the performance of PostgreSQL, SQLite, and DuckDB in executing quantitative queries by focusing on assessing query execution time, output quality, resource utilization, scalability, and security across these databases, aiming to identify the most efficient system for processing large datasets. The methodology involves the following steps:

Category	Subcategory	Description
Data Collection	Data Sources	Data was obtained from two CSV files containing appliance metadata and reviews, including attributes such as price, categories, and descriptions.
	Database Setup	The data was imported into PostgreSQL, SQLite, and DuckDB. PostgreSQL and SQLite were configured for traditional SQL queries, while DuckDB was utilized for in-process analytical querying.
Query Design	Objective	The primary goal was to compute the average price of water filters and evaluate performance metrics.
	Query Structure	A standardized SQL query was constructed to calculate the average price of water filters. The query was adjusted for compatibility with PostgreSQL, SQLite, and DuckDB based on each database's syntax and capabilities.
Performance Measurement	Query Execution Time	The duration from query initiation to result retrieval was recorded for each database to assess performance efficiency.
	Latency	Evaluated based on query execution time.
Quality of Output	Accuracy Verification	Results were compared to verify the accuracy of the average price calculations. Discrepancies were analyzed to understand differences in numeric data handling and type conversions.

Resource Utilization	Resource Consumption	Memory and CPU usage were assessed during query execution.
Scalability and Cloud Support	Scalability Assessment	Evaluated each database's ability to manage large datasets and scale horizontally or vertically.
	Cloud Support	Reviewed compatibility and support for cloud deployment based on available documentation and managed services.
Cost and Security Analysis	Cost Evaluation	Analyzed operational and cloud deployment costs for each database.
	Security Features	Reviewed security features such as encryption and access controls.
Results and Recommendations	Data Analysis	Compared performance metrics to identify strengths and weaknesses.
	Recommendations	Provided suggestions for the most suitable database based on performance, scalability, and cost considerations.

TABLE 2.1 Quantitative Queries Research Methodology

2.2 Fuzzy Queries:

Fuzzy queries were tested using Elasticsearch and Pinecone to evaluate their performance in handling approximate search requirements. The methodology includes:

Category	Subcategory	Description
Data Collection	Data Sources	The same appliance metadata and reviews data were used.
	Database Setup	Elasticsearch was set up using a Docker image for a consistent environment, with data indexed via Python's Pandas library. The database was accessed through the Python client library at http://localhost:9200 for efficient querying. Pinecone utilized OpenAI's 'text-embedding-ada-002' model to generate vector embeddings. With a token limit of 1,000,000 tokens per minute and 10,431,473 tokens from the review text, data was ingested in batches to optimize performance.
Query Design	Objective	To assess the performance of fuzzy queries in handling variations in user input and to compare document-based and vector-based approaches.

	Query Structure	Constructed fuzzy queries using Elasticsearch's Boolean queries, which handle approximate matches. In Pinecone, vector embeddings were used with Euclidean distance to evaluate similarity based on the review text.
Performance Measurement	Query Execution Time	Recorded the time taken to execute fuzzy queries on both Elasticsearch and Pinecone.
	Latency	Evaluated latency in retrieving results.
Quality of Output	Accuracy and Relevance	Manually evaluated the accuracy and relevance by analyzing the top 10 records returned from each system. The aim was to compare the performance of document-based and vector-based databases in providing relevant search results.
Resource Utilization and Scalability	Resource Consumption	Analyzed memory and CPU usage during query execution.
	Scalability	Assessed each system's ability to handle large volumes of fuzzy queries and scale efficiently, including evaluating the impact of batch ingestion and dimensional adjustments in Pinecone.
Cost and Security Analysis	Cost Evaluation	Reviewed the operational and cloud costs associated with running fuzzy queries on both Elasticsearch and Pinecone.
	Security Features	Assessed relevant security measures for handling fuzzy queries.

TABLE 2.2 Fuzzy Queries Research Methodology

2.3 Hybrid Queries

Hybrid queries were tested using Milvus and MongoDB to evaluate their performance in handling approximate search requirements. The methodology includes:

Category	Subcategory	Description
Data Collection	Data Sources	Data was obtained from two CSV files containing appliance metadata and reviews, including attributes such as price, categories, and descriptions

	Database Setup	This research uses two databases: Milvus for vector-based search and MongoDB for hybrid search. Milvus handles high-dimensional vector embeddings, while MongoDB supports vector search with Lucene-based indexing for hybrid functionality.
Query Design	Objective	The main objective is to evaluate hybrid search performance by combining vector and keyword searches to retrieve customer reviews that match criteria like "Easy to use" within a specific time frame.
	Query Structure	Queries combine vector-based similarity search with traditional keyword matching, embedding text data into vectors for similarity comparisons and using keyword filters to refine results.
Performance Measurement	Query Execution Time	Performance is measured by query execution time. Milvus, optimized for vector search, is expected to have faster execution, while MongoDB's speed may vary with query complexity.
	Latency	Latency is crucial for real-time applications. This research compares Milvus and MongoDB to see which offers more responsive query execution for hybrid search under similar conditions.
Quality of Output	Accuracy Verification	The accuracy of search results is assessed by comparing similarity scores and relevance. Higher scores and more relevant matches indicate better performance. MongoDB is expected to show slightly higher relevance due to its integrated capabilities, while Milvus focuses on vector-based similarity.
Resource Utilization		Resource utilization measures each database's computational efficiency. Milvus, optimized for vector search, may use more CPU/GPU resources. MongoDB, being general-purpose, is expected to balance resource use across various queries.
Scalability and Cloud Support	Scalability Assessment	Both databases are assessed for scalability with growing data volumes. Milvus excels in large-scale vector searches, while MongoDB offers strong general-purpose scalability, making it suitable for diverse applications.

	Cloud Support	Cloud support is assessed by how each database integrates with cloud environments. MongoDB offers seamless cloud integration through its service, MongoDB Atlas. Milvus is also cloud-friendly, especially in Kubernetes deployments.
Cost and Security Analysis	Cost Evaluation	The cost analysis looks at both operational and infrastructure costs. Milvus might be more cost-effective for specialized vector searches, while MongoDB could be more expensive depending on query complexity and cloud service use.
	Security Features	Security is crucial. MongoDB offers comprehensive features like encryption, access control, and cloud-based enhancements. Milvus has standard features but may need extra configurations to match MongoDB's level.
Results and Recommendations	Data Analysis	Milvus excels in vector search with low latency and efficient resource use. However, MongoDB is more versatile, offering better output quality, broader cloud support, and stronger security features.

TABLE 2.3 Hybrid Queries Research Methodology

Comparative Analysis/Results

3.1 Quantitative Queries

Query: *Find average price of water filters*

Criteria	PostgreSQL	SQLite	DuckDB
Query Execution Time	2.079 seconds	0.04 seconds	0.0084 seconds
Average Price for Water Filters	47.66381	44.16	47.5499
Latency	Higher latency due to server-based architecture	Very low latency due to in-process execution	Extremely low latency, optimized for analytics

Quality of Output	Excellent, with support for complex queries	Good for basic to moderate needs, lacks some advanced features	High quality for analytical tasks and complex queries
Scalability	Good, with vertical and horizontal scaling options	Limited, best for single-user or small-scale scenarios	Good for local analytics, not designed for distributed environments
Resource Utilization	Higher due to server-based nature and feature set	Very low, minimal resource consumption	Efficient for analytical queries, low overhead
Cloud Support	Excellent, with managed services available	Limited, not typically used in cloud environments	Growing, with increasing cloud support but primarily for local analytics
Cost	Free and open-source, cloud costs vary	Free and open-source, minimal operational costs	Free and open-source, no significant additional costs
Security	Strong security features including encryption and access controls	Basic security features, lacks advanced enterprise-level security	Basic security, suitable for local environments
Architecture of DB	Client-server architecture, robust feature set and scalability	Embedded, serverless, self-contained	In-process OLAP database, optimized for in-memory analytics

TABLE 3.1. PostgreSQL vs DuckDB vs SQLite

Observations on Average Price Discrepancy

SQLite shows a slightly lower average price (44.16) for water filters compared to PostgreSQL (47.66381) and DuckDB (47.5499). (Check Appendix A Fig (i-v))This difference arises because SQLite uses dynamic typing and implicit type conversions, which can lead to variations, especially with numeric and text-based data. Additionally, SQLite's handling of NULL values and automatic type conversions may cause inconsistencies in

calculations. In contrast, PostgreSQL and DuckDB use explicit numeric types and more consistent NULL handling, resulting in more accurate and stable average prices.

From Table 3.1, we can see that DuckDB offers fast, efficient performance for local or embedded analytics. On the other hand, PostgreSQL is a robust, scalable solution with extensive cloud support, ideal for complex, secure, large-scale applications. SQLite is lightweight and simple, perfect for embedded use cases where advanced features aren't needed.

3.2 Fuzzy Queries:

Query: Retrieve all reviews mentioning "difficult to understand" and similar phrases.

For above queries top 10 records were tested in both Elasticsearch & Pinecone:

Criteria	Elasticsearch	Pinecone
Latency	0.23 seconds	0.84 seconds
Indexing Time	47.9 seconds	~5 hours
Storage Size	227 bytes	~813 MB
Quality of Output	Q3:10/10	Q3: 10/10
Scalability	Horizontal scaling with shards and nodes.	Designed for large-scale vector datasets, it automatically scales with data size.
Industry Use case	Wikipedia's full-text search.	Spotify's music recommendation system.
Resource Utilization	High CPU usage is required for text analysis and aggregations, significant RAM is needed for caching, and SSD storage is recommended for optimal performance.	Utilizes GPU for faster vector computations, optimized in-memory operations, and efficient vector compression techniques for storage.
Cloud Support	Elastic Cloud offers native Elasticsearch services on AWS, Google Cloud (GCP), and Azure.	Pinecone Cloud is fully managed and available on AWS, GCP, and Azure.

Cost	Elastic Cloud starts at \$95/month, AWS on-demand instances at \$0.10/hour, with self-hosted costs varying by infrastructure.	The pay-per-use model starts at ~\$0.00045 per GB/Hour, with a free version offering 2GB storage/5 indexes, plus extra costs for queries, data transfer, and OpenAI embeddings at \$0.010 per 1M tokens.
Security	Field and document-level security are protected by SSL/TLS encryption, with role-based access control and auditing support.	The system uses AES 256-bit encryption, offers IAM-style access controls with VPC isolation, and is SOC 2 Type 2 certified.
Architecture	A distributed search engine with RESTful API, ELK Stack integration, and multi-language support including Java, Python, .NET, and Go.	A serverless distributed vector database with a simple API (REST/gRPC) and native SDKs for Python, JavaScript, and Java.

TABLE 3.2. ElasticSearch vs Pinecone

Elasticsearch is best for traditional text search and efficient indexing due to its lower latency and compact storage, ideal for handling and retrieving textual data quickly.

Pinecone excels in advanced vector-based similarity searches and scalable, personalized recommendations, making it suitable for applications like tailored suggestions based on user reviews and preferences. Combining both technologies offers a robust solution for both efficient search and sophisticated personalization.

3.3 Hybrid Queries:

Query: *Get a list of those reviews that are similar to this text: "Easy to use" and got reviewed between Jan 2013 to August 2013.*

For above queries top 10 records were tested in both Milvus & MongoDB:

Criteria	Milvus	MongoDB
Query Execution Time	0.3239 seconds	0.6302 seconds

Throughput Time/Latency	Optimized for vector search, providing lower latency.	Generally slower for vector search compared to Milvus.
Quality of Output	Good quality output with accurate vector similarity.	Slightly higher similarity scores, potentially more relevant results.
Scalability	Highly scalable for large-scale vector searches.	Scalable for general-purpose use but less so for vector-specific tasks.
Resource Utilization	Efficient for vector operations, may require more computational resources.	Balanced resource usage but less optimized for vector searches.
Cloud Support	Cloud-friendly, with Kubernetes support.	Extensive cloud support, especially with MongoDB Atlas.
Cost	Potentially lower cost for vector search tasks.	Can be more expensive, depending on usage.
Security	Standard security features.	Comprehensive security features, including encryption and access control.
Architecture of DB	Specialized for high-performance vector search.	General-purpose with flexible schema support.

TABLE 3.3. Milvus vs MongoDB

Milvus is ideal for specialized, high-performance vector search tasks with lower latency and better scalability in this domain.

MongoDB offers broader capabilities with better cloud support, security, and general-purpose scalability, making it more versatile for varied use cases.

Recommendations

Database Name	Best For	Recommendation
DuckDB	Fast in-memory analytics, high-speed data processing, local or embedded scenarios.	Use DuckDB for efficient analytical queries on large datasets within a single environment, ideal for data science and performance-critical tasks.
PostgreSQL	Comprehensive feature set, high-security environments, scalable applications.	Choose PostgreSQL for applications needing complex queries, high reliability, robust security, and extensive cloud support.
SQLite	Lightweight, embedded, or single-user applications.	Choose SQLite for lightweight, low-overhead database needs in small to moderate applications where simplicity and minimal resource use are key.
Elasticsearch	Elasticsearch is ideal for fast, efficient text search and indexing, perfect for large datasets like Wikipedia.	Use Elasticsearch for fast indexing and searching of large text volumes with minimal storage, ideal for speed-critical applications handling both structured and unstructured data.
Pinecone	Pinecone excels in scalable vector-based similarity searches, ideal for large-scale personalized recommendation systems.	Choose Pinecone for high-quality, scalable vector-based recommendations, especially in AI-driven applications requiring precise similarity searches.
Milvus	Milvus is ideal for high-performance vector search in AI/ML applications involving large datasets and efficient similarity search.	Milvus is ideal for vector-based similarity searches, offering top performance and scalability for large-scale AI/ML tasks.
MongoDB	MongoDB excels in hybrid search scenarios, ideal for enterprise applications needing robust cloud support, security, and versatility.	MongoDB is ideal for hybrid search, offering a balanced mix of performance, cost, and flexibility, making it versatile for various applications.

TABLE 4 Recommendations

Conclusion

This study reveals the diverse strengths of database systems for various query types. It emphasizes matching database capabilities with organizational needs as data complexity

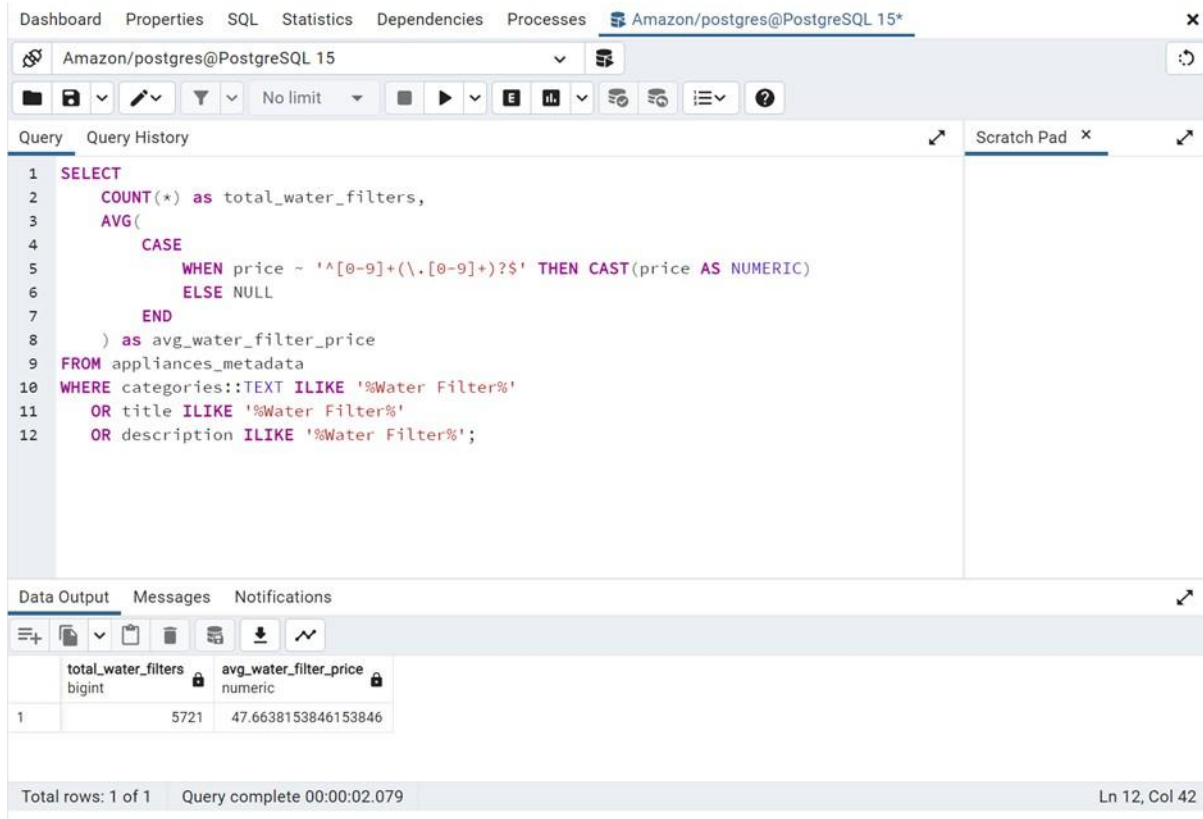
grows. DuckDB excels in fast analytics, while PostgreSQL suits complex applications. For fuzzy searches, Elasticsearch provides efficient text indexing, and Pinecone shines in vector-based recommendations. In hybrid queries, Milvus outperforms in vector search speed, while MongoDB offers versatility. Selecting the right database is crucial for efficient data management and querying. Future research should explore emerging technologies and AI integration to address evolving data challenges in our increasingly data-driven world.

References

- Graft, J., N. H. Nødtvedt, A. D. Pimentel, and M. W. Fagerland. "Computational Efficiency of Statistical Analyses: A Comparison of R, Python, Julia and SQLite." *PLoS ONE* 16, no. 4 (2021): e0249216.
- Lahitani, Alfirna Rizqi, Adhistya Erna Permanasari, and Noor Akhmad Setiawan. Cosine Similarity to Determine Similarity Measure: Study Case in Online Essay Assessment. Department of Electrical Engineering and Information Technology, Faculty of Engineering, Universitas Gadjah Mada. Accessed August 23, 2024.
- Levenshtein, Vladimir I. "Binary Codes Capable of Correcting Deletions, Insertions, and Reversals." *Soviet Physics Doklady* 10, no. 8 (1966): 707-710.
- Özsu, M. T. "A Survey of RDF Data Management Systems." *Frontiers of Computer Science* 14, no. 2 (2020): 217-252.
- Pan, James Jie, Jianguo Wang, and Guoliang Li. 2024. "Survey of Vector Database Management Systems." *The VLDB Journal* 33:1591–1615. <https://doi.org/10.1007/s00778-024-00864-x>.
- Raasveldt, M., and H. Mühleisen. "DuckDB: An Embeddable Analytical Database." *Proceedings of the 2019 International Conference on Management of Data* (2019): 1981-1984.
- Reinsel, David, John Gantz, and John Rydning. *Data Age 2025: The Evolution of Data to Life-Critical. Don't Focus on Big Data; Focus on the Data That's Big*. IDC White Paper, sponsored by Seagate, April 2017. Accessed August 23, 2024. <https://www.seagate.com/www-content/our-story/trends/files/Seagate-WP-DataAge2025-March-2017.pdf>.
- Sawhney, Navneet, Bose, Anirban and Paranthaman, Eswarankarthik. "A Report on Comparative Analysis of LLMs on Different Categories of Queries: Phase 1." Mid term project, MS DSP 420- Database Systems, Professor Abid Ali, July 27, 2024.
- Shaik, Subhani, and Nallamotheu Naga Malleswara Rao. 2017. "A Review of Elasticsearch: Performance Metrics and Challenges." *International Journal on Recent and Innovation Trends in Computing and Communication* 5, no. 11: 222-229. Available at <http://www.ijritcc.org>.
- Wallen, Jack. "Comparing DBMSs: The 8 Most Popular Databases." *Linode*, February 25, 2022. Accessed August 23, 2024. <https://www.linode.com/docs/guides/list-of-databases/>.
- Xie, Xingrui, Wenzhe Hou, Han Liu, and Hongbin Huang. 2023. "A Brief Survey of Vector Databases." *Proceedings of the 2023 9th International Conference on Big Data and Information Analytics (BigDIA)*. IEEE.

Appendix A

Quantitative Query Codes Screenshots



The screenshot shows a PostgreSQL query editor interface. The top bar includes tabs for Dashboard, Properties, SQL, Statistics, Dependencies, and Processes. The main query editor displays the following SQL query:

```

1 SELECT
2     COUNT(*) as total_water_filters,
3     AVG(
4         CASE
5             WHEN price ~ '^[0-9]+(\.[0-9]+)?$' THEN CAST(price AS NUMERIC)
6             ELSE NULL
7         END
8     ) as avg_water_filter_price
9 FROM appliances_metadata
10 WHERE categories::TEXT ILIKE '%Water Filter%'
11        OR title ILIKE '%Water Filter%'
12        OR description ILIKE '%Water Filter%';

```

Below the query editor, the Data Output tab shows the results of the query:

	total_water_filters bigint	avg_water_filter_price numeric
1	5721	47.6638153846153846

The bottom status bar indicates: Total rows: 1 of 1, Query complete 00:00:02.079, Ln 12, Col 42.

Figure i. Code for Quant Query using Postgres

```

1: import duckdb
import time

# Connect to DuckDB
con = duckdb.connect('my_database.duckdb')

# Load data into DuckDB
# con.execute("""
# CREATE TABLE appliances_metadata AS SELECT * FROM read_csv_auto('C:/Users/navne/Downloads/Final_Project_Phase_1/Final_Project_Phase_1/Amazon/Amazon_Applianc
# CREATE TABLE appliances_reviews AS SELECT * FROM read_csv_auto('C:/Users/navne/Downloads/Final_Project_Phase_1/Final_Project_Phase_1/Amazon/Amazon_Appliance
# """)

# Query: Find average price of water filters
start_time = time.time()
result_avg_price = con.execute("""

-- Average price of water filters
SELECT AVG(CAST(price AS FLOAT)) as avg_water_filter_price
FROM appliances_metadata_view
WHERE categories LIKE '%Water Filter%'
       OR title LIKE '%Water Filter%'
       OR description LIKE '%Water Filter%';

""").fetchall()
end_time = time.time()
print(f"Query execution time for average price: {end_time - start_time} seconds")
print(result_avg_price)

|
4

```

Query execution time for average price: 0.014490365982055664 seconds
[(47.5499279821938,)]

Figure ii. Code for Quant Query using DuckDB

```

25]: import sqlite3
import csv
import time

# Connect to SQLite database (creates a new file if it doesn't exist)
conn = sqlite3.connect('amazon_appliances.db')
cursor = conn.cursor()

# Create tables
cursor.execute('''
CREATE TABLE IF NOT EXISTS appliances_metadata (
    asin TEXT PRIMARY KEY,
    price TEXT,
    imageUrl TEXT,
    description TEXT,
    categories TEXT,
    title TEXT,
    brand TEXT,
    related TEXT,
    salesRank TEXT
)
''')

cursor.execute('''
CREATE TABLE IF NOT EXISTS appliances_reviews (
    reviewerID TEXT,
    asin TEXT,
    reviewerName TEXT,
    helpful TEXT,
    reviewText TEXT,
    overall REAL,
    summary TEXT,
    unixReviewTime INTEGER,
    reviewTime TEXT,
    PRIMARY KEY (reviewerID, asin)
)
''')

```

```

# Function to import CSV data
def import_csv(file_path, table_name):
    start_time = time.time()
    with open(file_path, 'r', encoding='utf-8') as csvfile:
        csv_reader = csv.reader(csvfile)
        headers = next(csv_reader)
        placeholders = ','.join(['?' for _ in headers])
        insert_query = f"INSERT OR REPLACE INTO {table_name} VALUES ({placeholders})"

        cursor.executemany(insert_query, csv_reader)

    conn.commit()
    end_time = time.time()
    print(f"Imported {table_name} in {end_time - start_time:.2f} seconds")

# Import data
import_csv('C:/Users/navne/Downloads/Final_Project_Phase_1/Final_Project_Phase_1/Amazon/Amazon_Appliances_Metadata.csv', 'appliances_metadata')
import_csv('C:/Users/navne/Downloads/Final_Project_Phase_1/Final_Project_Phase_1/Amazon/Amazon_Appliances_Reviews.csv', 'appliances_reviews')

# Verify data import
cursor.execute("SELECT COUNT(*) FROM appliances_metadata")
print("Metadata count:", cursor.fetchone()[0])

cursor.execute("SELECT COUNT(*) FROM appliances_reviews")
print("Reviews count:", cursor.fetchone()[0])

start_time = time.time()
# Example queries

# Average rating for water filters
cursor.execute("""
SELECT
    COUNT(*) as total_water_filters,
    AVG(
        CASE
            WHEN price GLOB '[0-9]*.[0-9]*' THEN CAST(price AS REAL)
            ELSE NULL
        END
    ) as avg_water_filter_price
FROM appliances_metadata
WHERE categories LIKE '%Water Filter%'
OR title LIKE '%Water Filter%'
OR description LIKE '%Water Filter%';
""")

```

```

result = cursor.fetchone()
end_time = time.time()

avg_price = result[1]

print(f"Average price for water filters: {avg_price:.2f}")
print(f"Query execution time: {end_time - start_time:.2f} seconds")

```

```

Imported appliances_metadata in 0.69 seconds
Imported appliances_reviews in 8.10 seconds
Metadata count: 11656
Reviews count: 143685
Average price for water filters: 44.16
Query execution time: 0.04 seconds

```

Figure iii,iv and v. Code for Quant Query using SQLite

Fuzzy Query Codes Screenshots

Elasticsearch

```

query_3 = {
  "bool": {
    "should": [
      {
        "match_phrase": {
          "reviewText": {
            "query": "difficult to understand",
            "boost": 2.0
          }
        }
      },
      {
        "match_phrase": {
          "reviewText": {
            "query": "hard to comprehend",
            "boost": 1.5
          }
        }
      },
      {
        "match_phrase": {
          "reviewText": "confusing"
        }
      }
    ],
    "minimum_should_match": 1
  }
}

```

Figure vi. Code for Elastic Search.

```

def run_scroll_query(es_client, index_name, query, scroll='2m', size=1000):
    start_time = time.time()

    try:
        response = es_client.search(
            index=index_name,
            body={
                "query": query,
                "size": size,
                "_source": ["reviewText", "reviewerID", "asin"],
                "track_total_hits": True
            },
            scroll=scroll
        )

        scroll_id = response['_scroll_id']
        hits = response['hits']['hits']

        total_docs = len(hits)
        all_hits = hits.copy()

        while len(hits) > 0:
            response = es_client.scroll(scroll_id=scroll_id, scroll=scroll)
            scroll_id = response['_scroll_id']
            hits = response['hits']['hits']

            total_docs += len(hits)
            all_hits.extend(hits)

        es_client.clear_scroll(scroll_id=scroll_id)

        end_time = time.time()
        latency = end_time - start_time
        throughput = total_docs / latency if latency > 0 else 0

        return total_docs, latency, throughput, all_hits

```

Figure vi. Code for Elastic Search.

```

except Exception as e:
    print(f"An error occurred: {e}")
    return 0, 0, 0, []

# Data retrieval and processing of data:
total_docs_3, latency_3, throughput_3, all_hits_3 = run_scroll_query(es, index_name, query_3)

# Datframe conversion:
df_3 = pd.DataFrame([
    {**hit['_source'], 'score': hit.get('_score', None)}
    for hit in all_hits_3
])

# Performance metrics:
print("Overall Performance:")
print("-----")
print(f"Latency: {latency_3:.2f} seconds")
print(f"Throughput: {throughput_3:.2f} documents/second")

# Display the top 10 searches:
if not df_3.empty:
    top_matches_3 = df_3.sort_values(by='score', ascending=False).head(10)

    print("\nTop matches based on review text:")
    print("-----")

    for rank, (_, row) in enumerate(top_matches_3.iterrows(), start=1):
        print(f"Rank: {rank}")
        print(f"Score: {row.get('score', 'N/A')}")
        print(f"ReviewerID: {row.get('reviewerID', 'N/A')}")
        print(f"ASIN: {row.get('asin', 'N/A')}")
        print(f"Review Text: {row.get('reviewText', 'N/A')}")
        print()
else:
    print("No matches found.")

```

Figure vii. Code for Elastic Search.

Overall Performance:

 Latency: 0.21 seconds
 Throughput: 879.55 documents/second

Top matches based on review text:

Rank: 1
 Score: 27.246609
 ReviewerID: A15U4KORNHPCXH
 ASIN: B004H3XWCO
 Review Text: the instructions are difficult to understand for the average person and i will have to have a service call anyway.

Rank: 2
 Score: 27.246609
 ReviewerID: A1TGMF7WCRXB7
 ASIN: B00E0XT3G
 Review Text: it is nice but the instructions are in chinese a little difficult to understand without the english instructions i guessmarco

Rank: 3
 Score: 26.751848
 ReviewerID: A30ZADVSR9JUYE
 ASIN: B001JEOIFY
 Review Text: a little difficult to understand at first, but once you get the hang of it, a nice addition to your tool collection.

Rank: 4
 Score: 9.494004
 ReviewerID: AHE8EFUW0T009
 ASIN: B00DM81IOQ
 Review Text: kit was slightly confusing, didn't quite match my refrigerator but i got it to work

Rank: 5
 Score: 9.139553
 ReviewerID: A6X9Z7EPRVM6F
 ASIN: B003BIGDJO
 Review Text: replaced existing icemaker - directions were a bit confusing and it seemed we did it wrong but. . .we're making ice!

Figure viii. Elastic Search Output

Rank: 6
 Score: 9.055038
 ReviewerID: A1CDKODRUGCHLU
 ASIN: B001XW8KM4
 Review Text: my husband said the directions were a bit confusing. otherwise, the product was fine and it was easy to install.

Rank: 7
 Score: 8.890612
 ReviewerID: A3HP6FV6D4HHSR
 ASIN: B0014X7B54
 Review Text: very easy to install and somewhat quieter than insinkerator.the instructions were confusing,but could be figured out. i would recommend this brand.

Rank: 8
 Score: 8.890612
 ReviewerID: AOHJW6CTC5STN
 ASIN: B0050KKM62
 Review Text: the directions are a bit confusing, i couldn't determine which tab the jumper went on. it went down to trial and error.

Rank: 9
 Score: 8.654871
 ReviewerID: A392365U18C7ZW
 ASIN: B004XLE3RI
 Review Text: i did not receive the bottom meat pan/crisper as the picture shows here but the other vegetable crisper drawer above it. .confusing. any tips?

Rank: 10
 Score: 8.654871
 ReviewerID: AEUER3OAXTT5Q
 ASIN: B006R8A28E
 Review Text: i would have given it a 5 star but the enclosed wire w/o instructions was confusing. as it turned out the wire was notnecessary.

Figure ix. Elastic Search Output

Pinecone:

```
# OpenAI client initialisation:
OPEN_AI_API_KEY = '*****'
openai_client = OpenAI(api_key=OPEN_AI_API_KEY)

# Pinecone client instance initialisation:
PINECONE_API_KEY = '****'
pc = Pinecone(api_key=PINECONE_API_KEY)

# Top 10 records:
TOP_K = 10

def main():
    query = "difficult to understand OR hard to follow OR confusing OR unclear OR hard to comprehend" # Query-3
    results, latency = search_reviews(query, index, TOP_K)

    if results and 'matches' in results:
        num_documents = len(results['matches'])
        throughput = num_documents / latency if latency > 0 else 0
        print("Overall Performance:")
        print("-----")
        print(f"Latency: {latency:.2f} seconds")
        print(f"Throughput: {throughput:.2f} documents/second")
        print()
        print(f"Top {TOP_K} matches for '{query}':")
        print("-----")

        for i, match in enumerate(results['matches'], 1):
            print(f"Rank: {i}")
            print(f"Score: {match['score']:.6f}")
            #print(f>Title: {match['metadata'].get('title', 'N/A')}")
            print(f"Author: {match['metadata'].get('author', 'N/A')}")
            print(f"ASIN: {match['metadata'].get('asin', 'N/A')}")
            print(f"Review Text: {match['metadata'].get('reviewText', 'N/A')}")
            #print(f>Description: {match['metadata'].get('description', 'N/A')}")
            #print(f"Summary: {match['metadata'].get('summary', 'N/A')}")
            print()
```

Figure x. Code for Pinecone

```
else:
    print("No relevant reviews found.")

# Main function:
main()
```

Figure xi. Code for Pinecone

```
Overall Performance:
-----
Latency: 0.86 seconds
Throughput: 11.56 documents/second

Top 10 matches for 'difficult to understand OR hard to follow OR confusing OR unclear OR hard to comprehend':
-----
Rank: 1
Score: 0.359810
Author: N/A
ASIN: B001JE0IFY
Review Text: a little difficult to understand at first, but once you get the hang of it, a nice addition to your tool collection.

Rank: 2
Score: 0.372300
Author: N/A
ASIN: B00GHXU3VA
Review Text: instructions not easy to understand. it is up and it works, thought it would be a little better quality.

Rank: 3
Score: 0.374198
Author: N/A
ASIN: B004H3XWCO
Review Text: the instructions are difficult to understand for the average person and i will have to have a service call anyway.

Rank: 4
Score: 0.392511
Author: N/A
ASIN: B001EJD7NI
Review Text: could not figure out their use.

Rank: 5
Score: 0.394586
Author: N/A
ASIN: B004H3XWCO
Review Text: the product works really well, completely solved my problem. instructions were very hard to understand and follow. we had to go on-line to get better instructions that made sense.
```

Figure xii. Pinecone Output

Rank: 6
Score: 0.400244
Author: N/A
ASIN: B003MU95X8
Review Text: difficult to read in normal room light. you need to be right on top of the indoor unit to read.

Rank: 7
Score: 0.401624
Author: N/A
ASIN: B001737LGU
Review Text: well made product, except it had no instructions in the package despite the package stating "easy to follow instructions enclosed" right on the front. . .

Rank: 8
Score: 0.403811
Author: N/A
ASIN: B00E4Q006U
Review Text: not telescoping. made the entire process really difficult. the specifications where not clear. wasted a lot of our time and money.

Rank: 9
Score: 0.406334
Author: N/A
ASIN: B0037MBG5Q
Review Text: it took me a tad of time to understand the set-up directions--but then again--i'm not the brightest lamp inna room.

Rank: 10
Score: 0.407065
Author: N/A
ASIN: B000UN2DTE
Review Text: this one did not have the usual easy to read instructions on it anywhere. it's impossible to remember after 6 months the proper way to install.

Figure xiii. Pinecone Output

Hybrid Query Codes Screenshots

Milvus

```
In [51]: connections.connect(host=HOST, port=PORT)

In [53]: # Remove collection if it already exists
if utility.has_collection(COLLECTION_NAME):
    utility.drop_collection(COLLECTION_NAME)

In [55]: # Create collection which includes the id, title, and embedding.
fields = [
    FieldSchema(name='reviewerID', dtype=DataType.VARCHAR, max_length=64000, is_primary=True),
    FieldSchema(name='asin', dtype=DataType.VARCHAR, max_length=64000),
    FieldSchema(name='reviewerName', dtype=DataType.VARCHAR, max_length=64000),
    FieldSchema(name='helpful', dtype=DataType.VARCHAR, max_length=64000),
    FieldSchema(name='reviewText', dtype=DataType.VARCHAR, max_length=64000),
    FieldSchema(name='overall', dtype=DataType.FLOAT),
    FieldSchema(name='summary', dtype=DataType.VARCHAR, max_length=64000),
    FieldSchema(name='unixReviewTime', dtype=DataType.INT64),
    FieldSchema(name='reviewTime', dtype=DataType.VARCHAR, max_length=64000),
    FieldSchema(name='embedding', dtype=DataType.FLOAT_VECTOR, dim=DIMENSION)
]

schema = CollectionSchema(fields=fields)
collection = Collection(name=COLLECTION_NAME, schema=schema)

In [57]: # Create the index on the collection and load it.
collection.create_index(field_name="embedding", index_params=INDEX_PARAM)
collection.load()

In [58]: from openai import OpenAI
client = OpenAI()

In [59]: # Simple function that converts the texts to embeddings
def embed(texts):
    embeddings = client.embeddings.create(
        input=texts,
        model=OPENAI_ENGINE
    )
    return [x.embedding for x in embeddings.data]

In [60]: def convert_datetime_to_unix_time(datetime_obj):
date_format = datetime.datetime.strptime(datetime_obj, "%Y-%m-%dT%H:%M:%SZ")
unix_time = datetime.datetime.timestamp(date_format)

return int(unix_time)

In [62]: list_of_all_issues_fetched = appliances_reviews_df.values.tolist()

In [63]: list_of_all_issues_fetched[0:4]
```

Figure xiv. Code for Hybrid Query setup using Milvus

```

In [65]: from tqdm import tqdm

data = [
    [], # reviewerID
    [], # asin
    [], # reviewerName
    [], # helpful
    [], # reviewText
    [], # overall
    [], # summary
    [], # unixReviewTime
    [], # reviewTime
]

# Embed and insert in batches
for i in tqdm(range(0, 1900)):
    #for i in tqdm(range(0, 4)):
        data[0].append(str(list_of_all_issues_fetched[i][0])) # reviewerID
        data[1].append(str(list_of_all_issues_fetched[i][1])) # asin
        data[2].append(str(list_of_all_issues_fetched[i][2]) if list_of_all_issues_fetched[i][2] else "") # reviewerName
        data[3].append(str(list_of_all_issues_fetched[i][3])) # helpful
        data[4].append(str(list_of_all_issues_fetched[i][4])) # reviewText
        data[5].append(float(list_of_all_issues_fetched[i][5])) # overall
        data[6].append(str(list_of_all_issues_fetched[i][6])) # summary
        data[7].append(int(list_of_all_issues_fetched[i][7])) # unixReviewTime
        data[8].append(str(list_of_all_issues_fetched[i][8])) # reviewTime

        if len(data[0]) % BATCH_SIZE == 0:
            data.append(embed(data[4]))
            collection.insert(data)
            data = [[],[],[],[],[],[],[],[],[]]

# Embed and insert the remainder
if len(data[0]) != 0:
    data.append(embed(data[4]))
    collection.insert(data)
    data = [[],[],[],[],[],[],[],[],[]]

100% ██████████ 1900/1900 [00:04<00:00, 418.89it/s]

```

Figure xv. Code for data load into Milvus

```

In [69]: # Flush to ensure data is persisted
collection.flush()

In [124]: import time
# Helper Function
# Filtered Search Function

# Adjust the top_k parameter value to whatever the number of issues you want to inspect/print
# Please note that you might get thousands of issues back
# For unit-testing purposes, inspect/print few of these issues
# set top_k = 10 initially, and later you could change that to 100, 1000, etc.

def ask_milvus(query, top_k = 10):
    text, expr = query

    # Start timing
    start_time = time.time()

    results = collection.search(embed(text), anns_field='embedding', expr=expr, param=QUERY_PARAM, limit = top_k,
                                output_fields = ['reviewerID', 'asin',
                                                  'reviewerName', 'reviewText',
                                                  'overall', 'summary',
                                                  'reviewTime'])

    # End timing
    end_time = time.time()

    # Calculate the processing time
    processing_time = end_time - start_time

    for i, hit in enumerate(results):
        print(f'\nShowing Top {top_k} Results for query "{text}":\n')
        for j, hits in enumerate(hit):
            print('\t' + 'Rank:', j + 1, '| Score:', hits.score)
            print('\t\t' + 'reviewerID:', hits.entity.get('reviewerID'))
            print('\t\t' + 'asin:', hits.entity.get('asin'))
            print('\t\t' + 'reviewerName:', hits.entity.get('reviewerName'))
            print('\t\t' + 'reviewText:', hits.entity.get('reviewText'))
            print('\t\t' + 'overall:', hits.entity.get('overall'))
            print('\t\t' + 'summary:', hits.entity.get('summary'))
            print('\t\t' + 'reviewTime:', hits.entity.get('reviewTime'))
            print("\n")

    # Print the processing time
    print(f"Processing Time: {processing_time:.4f} seconds")

<>:32: SyntaxWarning: invalid escape sequence '\S'
<>:32: SyntaxWarning: invalid escape sequence '\S'
C:\Users\bonjo\AppData\Local\Temp\ipykernel_3120\1351864662.py:32: SyntaxWarning: invalid escape sequence '\S'
print(f'\nShowing Top {top_k} Results for query "{text}":\n')

```

Figure xvi. Code for search query function in Milvus

```

# Convert date to unix-time
unix_time_start = convert_datetime_to_unix_time("2013-01-01T00:00:00Z")
unix_time_end = convert_datetime_to_unix_time("2024-08-31T23:59:00Z")

# Create your filter/expression for Milvus
#expr = 'unixReviewTime > ' + str(unix_time)
expr = f'unixReviewTime >= {unix_time_start} and unixReviewTime <= {unix_time_end}'
# Put together your query
query = ('Easy to use', expr)

# Send your query to Milvus using the following helper function
ask_milvus(query)

# Inspect (Score, title, and relevance to the repo) for every issue listed in the output/results:

```

Figure xvii. Code for executing the search query function in Milvus

\Showing Top 10 Results for query "Easy to use":

Rank: 1	Score: 0.178335294274643 reviewerID: A3PQ43267RPV91 asin: B000040J30 reviewerName: Terry N. reviewText: easy to install overall: 5.0 summary: Five Stars reviewTime: 07 10, 2014
Rank: 2	Score: 0.261328950465027 reviewerID: A3Q3A4818970EY asin: B000056JFC reviewerName: Eugene Garrippy reviewText: This is my first experience with the product, and I found it easy to use. I would not hesitate to recomend it overall: 5.0 summary: Easy to use reviewTime: 04 11, 2013
Rank: 3	Score: 0.273124635219574 reviewerID: A26C18PLV83K13 asin: B0000418K4 reviewerName: S. Nisang reviewText: This item is easy to use. There is no use of batteries and does not have to be programed to function. It also shows the range that works well when setting temperature and humidity in the home. overall: 5.0 summary: Quick view reviewTime: 01 10, 2013
Rank: 4	Score: 0.2841043770913263 reviewerID: A3T0W2K3Y96ZB asin: B00004XSF0 reviewerName: Brenda reviewText: Met our expectations. Very practical. Simple to use. Fits perfectly. So simple to use it is hard to write a review. overall: 5.0 summary: works great reviewTime: 01 10, 2013
Rank: 5	Score: 0.29545063278018035 reviewerID: A13K5G2NAB55M asin: B000040J30 reviewerName: D. Beach reviewText: Was easy install, exactly like my old one. overall: 5.0 summary: Five Stars reviewTime: 07 7, 2014
Rank: 6	Score: 0.2963576318033406 reviewerID: A1ABY081E1E12 asin: B00004VUN0 reviewerName: E.K. reviewText: very easy to work with and install. wires feel very secure and it comes with all hardware needed to install overall: 5.0 summary: Good outlet reviewTime: 01 10, 2013
Rank: 7	Score: 0.306377512840271 reviewerID: A2CV82D8B080UM asin: B00005AUX0 reviewerName: Jennifer williams reviewText: Very easy to use. I tested my water and right away knew if it was safe. I felt confident of the results. This was a lot cheaper than a private water testing company. overall: 5.0 summary: Good buy reviewTime: 10 10, 2013
Rank: 8	Score: 0.30870044261550093 reviewerID: A23609QAGH1OE asin: B00005AUX0 reviewerName: April Ebel reviewText: Product is just as it says. Directions are very helpful and easy to use. I would recommed for the price. overall: 5.0 summary: worked Great! reviewTime: 10 17, 2013
Rank: 9	Score: 0.31520854407310488 reviewerID: A2K08ECFF0F5D3 asin: B000040J30 reviewerName: SAL40r reviewText: Works like a charm overall: 5.0 summary: Five Stars reviewTime: 07 5, 2014
Rank: 10	Score: 0.3173400790695084 reviewerID: A3TFL6U6F0F0F asin: B000040J30 reviewerName: quick reviewText: works great, easy to install because the instructions lay out the work in a way anyone can follow diagram by diagram overall: 5.0 summary: no plumber required ! reviewTime: 04 4, 2014

Processing Time: 0.3299 seconds

Figure xviii. Output from Milvus search

MongoDB


```

In [63]: import requests
import datetime
import time
import os
import openai
import json

from datetime import date

In [114]: from pymongo.mongo_client import MongoClient
from pymongo.server_api import ServerApi
uri = "mongodb+srv://bonjoindia:Zm00qFU6dWk7ZZP5@cluster0.qa6hq3a.mongodb.net/?retryWrites=true&majority=&appName=Cluster0"
# Create a new client and connect to the server
mongo_client = MongoClient(uri, server_api=ServerApi('1'))
# Send a ping to confirm a successful connection
try:
    mongo_client.admin.command('ping')
    print("Pinged your deployment. You successfully connected to MongoDB!")
except Exception as e:
    print(e)

Pinged your deployment. You successfully connected to MongoDB!

In [67]: # GitHub Access Token

# Common API headers
headers = {
    "Accept": "application/vnd.github+json",

    "access_token": "ghp_yQwb4m0zQq9PsmTuzEVsnqaUzDVbxa0dnViv",
    "Git_Username": "anirbanbose83"
}

```

Figure xix. Code for Hybrid Query setup using MongoDB

```

In [ ]: from openai import OpenAI
client = OpenAI()

In [ ]: # Simple function that converts the texts to embeddings
def embed(texts):

    embeddings = client.embeddings.create(
        input=texts,
        model=OPENAI_ENGINE
    )

    return [x.embedding for x in embeddings.data]

In [ ]: def convert_datetime_to_unix_time(datetime_obj):
date_format = datetime.datetime.strptime(datetime_obj, "%Y-%m-%dT%H:%M:%SZ")
unix_time = datetime.datetime.timestamp(date_format)

return int(unix_time)

In [ ]: appliances_reviews_df_new = appliances_reviews_df.head(1990)

In [ ]: appliances_reviews_df_new["embedding"] = appliances_reviews_df_new['reviewText'].apply(embed)

C:\Users\bonjo\AppData\Local\Temp\ipykernel_3576\3397934495.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
appliances_reviews_df_new["embedding"] = appliances_reviews_df_new['reviewText'].apply(embed)

In [ ]: appliances_reviews_df_new

In [ ]: appliances_reviews_df_new

```

	reviewerID	asin	reviewerName	helpful	reviewText	overall	summary	unixReviewTime	reviewTime	embedding
0	A2THOWZ8UBJ8FF	0970408285	Steve	[0, 0]	Could have been longer though. well made and e...	4.0	Good fit	1387152000	12 16, 2013	[[-0.017633579671382904, -0.015432676300406456...
1	A24H084NFSTF5	7301113188	Maha Saqfalhait 'shopaholic j'	[0, 0]	I like these containers so much i have ordered...	5.0	I Love the Freezer storage line.	1236902400	03 13, 2009	[[-0.018233368173241615, -0.01791534572839737...
2	AXE83MK90ZEVZ	800002N7HY	Strom	[0, 0]	It works. no fires, etc. Why not 5 stars? Ho...	4.0	expectations achieved.	1389052800	01 7, 2014	[[-0.002737861592322588, 0.0066328574903309345...
3	A2J7X7ZIH2EWB1	800002NATH	NaN	[0, 0]	Fast shipping. Works great	5.0	Five Stars	1405814400	07 20, 2014	[[-0.02773250639438629, 0.016366302967071533, ...
4	AJQFN0FTZ7GOX	800002NATH	Barthbill	[1, 1]	What can I say? It is the usual Leviton high q...	5.0	good product at a good price.	1277164800	06 22, 2010	[[-0.014923588372766972, 0.018386593088507652, ...
...
1985	AN5WLTG3HPKCM	800005O64S	Maybelle Miller 'Gabcocklover'	[3, 4]	I have wanted a dishwasher for a long time. Un...	5.0	It's Great!	1222646400	09 29, 2008	[[-0.005219343584030867, -0.010897675529122353, ...
1986	A15X59BE180N1K	800005O64S	Media Enterprises, Inc.	[5, 5]	The unit started leaking and then quit functio...	1.0	LEAKED AND DIED	1265760000	02 10, 2010	[[-0.003302327822893858, 0.010290670208632946, ...
1987	A13MQ135487LBY	800005O64S	M. Filer	[8, 8]	I received the washer right on time, no proble...	1.0	Not worth it	1182816000	06 26, 2007	[[-0.021759038791060448, -0.006805486511439085, ...
1988	A1RWKR7L7VK28	800005O64S	Michael A. Kartchner 'kartvines'	[9, 9]	I have some concerns with this product. We hav...	3.0	Read all review before you buy	1131580800	11 10, 2005	[[-0.0030940892174839973, -0.0125398982457425594, ...
1989	A2P821GVMS9BXR	800005O64S	Mike	[0, 1]	This works great for just my wife and me. A li...	4.0	Works Great	1207785600	04 10, 2008	[[-0.006603363435715437, 0.0221431702375412, 0...

1990 rows x 10 columns

```

In [ ]: # Ingest data into MongoDB
db = mongo_client['assignment_group']
collection = db['Amazon_reviews_group_project']

In [ ]: documents = appliances_reviews_df_new.to_dict('records')
collection.insert_many(documents)

```

Figure xx. Code for loading embedded data into MongoDB

```

import time

def search_mongodb_vector(query_text, top_k=10, num_candidates=100):
    # Generate the embedding for the query text
    query_embedding = embed([query_text])[0] # Get the first (and only) embedding

    # Date range for filtering: January 2013 to August 2013
    start_date = datetime.datetime(2013, 1, 1)
    end_date = datetime.datetime(2024, 8, 31)

    # Start the timer
    start_time = time.time()

    # MongoDB vector search using the Lucene-based vector index
    results = collection.aggregate([
        {
            "$vectorSearch": {
                "index": "default", # Use the appropriate search index name if different
                "queryVector": query_embedding,
                "path": "embedding",
                "limit": top_k, # Limit the number of results to top_k
                "numCandidates": num_candidates # Number of candidates considered in the search
            }
        },
        {
            "$project": {
                "reviewerID": 1,
                "asin": 1,
                "reviewerName": 1,
                "reviewText": 1,
                "overall": 1,
                "summary": 1,
                "reviewTime": 1,
                "score": {"$meta": "vectorSearchScore"} # Include the similarity score
            }
        }
    ])

    filtered_results = []

```

Figure xxi. Code for search function in MongoDB


```

filtered_results = []

# Post-process filtering in Python
for result in results:
    review_time_str = result['reviewTime']
    review_time = datetime.datetime.strptime(review_time_str, "%m %d, %Y")

    if start_date <= review_time <= end_date:
        filtered_results.append(result)

# End the timer
end_time = time.time()

# Calculate the processing time
processing_time = end_time - start_time

# Print the filtered results
for i, result in enumerate(filtered_results):
    print(f"Rank: {i + 1} | Similarity Score: {result['score']}")
    print(f"Reviewer ID: {result['reviewerID']}")
    print(f"ASIN: {result['asin']}")
    print(f"Reviewer Name: {result['reviewerName']}")
    print(f"Review Text: {result['reviewText']}")
    print(f"Overall Rating: {result['overall']}")
    print(f"Summary: {result['summary']}")
    print(f"Review Time: {result['reviewTime']}")
    print("\n")

# Print the processing time
print(f"Processing Time: {processing_time:.4f} seconds")

```

```

# Example usage
search_mongodb_vector("Easy to use", top_k=10, num_candidates=1500)

```

Figure xxii. Code for triggering search function in MongoDB

```

Rank: 1 | Similarity Score: 0.954158184322785
Reviewer ID: A8PQA15w9P9Y5
ASIN: B000N4UJ10
Reviewer Name: Terry N.
Review Text: easy to install
Overall Rating: 5.0
Summary: Five Stars
Review Time: 07 10, 2014

Rank: 2 | Similarity Score: 0.9346748245813802
Reviewer ID: A3QQA816570EV
ASIN: B0000543FC
Reviewer Name: Eugene Carrieny
Review Text: This is my first experience with the product, and I found it easy to use. I would not hesitate to recommend it
Overall Rating: 5.0
Summary: Easy to use
Review Time: 08 11, 2013

Rank: 3 | Similarity Score: 0.9117748201958073
Reviewer ID: A25C1DPLV53R13
ASIN: B000N4UJ10
Reviewer Name: S. Nimog
Review Text: This item is easy to use. There is no use of batteries and does not have to be programed to function. It also shows the range that works well when setting temperature and humidity in the home.
Overall Rating: 5.0
Summary: Quick view
Review Time: 01 10, 2013

Rank: 4 | Similarity Score: 0.9289429187774658
Reviewer ID: A3DAV20X9V4Z8
ASIN: B000N4UJ10
Reviewer Name: Brenda
Review Text: Met our expectations. Very practical. Simple to use. Fits perfectly. So simple to use it is hard to write a review.
Overall Rating: 5.0
Summary: works great
Review Time: 01 10, 2013

Rank: 5 | Similarity Score: 0.92615958757446
Reviewer ID: A1JEG5U4M8550N
ASIN: B000N4UJ10
Reviewer Name: U. Baeth
Review Text: Was easy install, exactly like my old one.
Overall Rating: 5.0
Summary: Five Stars
Review Time: 07 7, 2014

Rank: 6 | Similarity Score: 0.9256492887225342
Reviewer ID: ALAB7081fFEI2
ASIN: B000N4UJ10
Reviewer Name: E.K.
Review Text: very easy to work with and install. wires feel very secure and it comes with all hardware needed to install
Overall Rating: 5.0
Summary: Good outlet
Review Time: 01 10, 2013

Rank: 7 | Similarity Score: 0.923469683861676
Reviewer ID: A2C8Z0R0M3M3M
ASIN: B000N4UJ10
Reviewer Name: Jennifer Williamsen
Review Text: Very easy to use. I tested my water and right away knew if it was safe. I felt confident of the results. This was a lot cheaper than a private water testing company.
Overall Rating: 5.0
Summary: Good buy
Review Time: 10 10, 2013

Rank: 8 | Similarity Score: 0.9227917104386455
Reviewer ID: A23R0X9C7H508
ASIN: B000N4UJ10
Reviewer Name: April Ebol
Review Text: Product is just as it says. Directions are very helpful and easy to use. I would recommend for the price.
Overall Rating: 5.0
Summary: Hurved Great!
Review Time: 10 17, 2013

Processing Time: 0.6382 seconds

```

Figure xxiii. Output from MongoDB